

C8: Atari Games with Deep Q Network (DQN)

▼ Table of Content

[Deep Q-Learning](#)

[Loss Function & Target Network](#)

[DQN Architecture](#)

[Tricks to Training a DQN](#)

[Experience Replay](#)

[Target Network](#)

[Rewards Clipping](#)

[Epsilon-Greedy Strategy](#)

[The Final Algorithm](#)

[Extra Tips and Tricks](#)

[Double DQN](#)

[Prioritized Experience Replay](#)


[Dueling Network Architecture](#)

<https://storage.googleapis.com/deepmind-media/dqn/DQNNaturePaper.pdf>

Playing Atari with Deep Reinforcement Learning

We present the first deep learning model to successfully learn control policies directly from high-dimensional sensory input using reinforcement learning. The model is a convolutional neural

 <https://www.arxiv-vanity.com/papers/1312.5602/>



Deep Q-Learning

- In Q-Learning, we estimate our Q-function through the construction of Q-table by finding the averaged return based on multiple iterations.
- Such exhaustive search across all possible state-action pairs for finding the optimal state soon faced a bottleneck when there are a large number of states and, in each state, there are a lot of actions to try.

- As such, Deep Q-Learning is introduced which uses NN with parameter/weights θ to approximate the Q-function with the goal that the approximated Q-function, $Q(s, a; \theta)$ is equiv/close to the optimal Q function, $Q^*(s, a)$.

$$Q(s, a; \theta) \approx Q^*(s, a)$$

Loss Function & Target Network

- Before we go into details of the actual network architecture, let us properly define the end-goal that we wish to achieve or in more specific terms, the loss function that we wish to minimize.
- Recall that our Q-network is nothing more but a neural network that can approximate the Q-function.
 - Thus, the number of output neuron for Q-network is just number of action available given the current state.

$$f : S \rightarrow A$$

- Recall that in Q-learning update rule, $r + \gamma \max Q(s', a)$ is the target value and $Q(s, a)$ is the predicted value and we tried to minimize the difference by learning a right policy.

$$Q(s, a) = Q(s, a) + \alpha[r + \gamma \max Q(s', a') - Q(s, a)]$$

- Hence, in DQN, our loss function is defined as the squared difference between the target and predicted value.

$$Loss = (y_i - Q(s, a; \theta))^2$$

- Where our target value, y_i is defined as follows

$$y_i = r + \gamma \operatorname{argmax}(Q(s', a'; \theta))$$

- Finally, the weights of Q-network is updated by minimizing the loss similar to how Q-learning works.

DQN Architecture

- Since the goal of a DQN is to approximate a Q-network, which learns from the rewards given by the environment after playing countless episodes of Atari games, a DQN is similar to a pure CNN with the first few layers being the convolutional network to learn from the pixel while the last few layers being the dense layer to approximate the Q-function
- The following are few characteristics of the CNN proposed in the paper and their reasoning behind the action:
 - Downsample 210x160 to 84x84 from RGB to Greyscaled input frame
 - To reduce the computational requirement for DQN
 - No pooling layer is used in the CNN
 - Retain positional information of pixels and features
 - Imagine a playing a Pong Game without the positional information of the ball, it is nearly impossible to do so.
 - Past four consecutive screen is passed in Forward Pass
 - Retain temporal information of the velocity of moving objects
 - Number of Output Neuron = Number of Action Space Possible
 - Get the Q-value for all actions available regardless of whether does a state allows

Tricks to Training a DQN

Experience Replay

- One more details when it comes to training a DQN is experience replay which is used to solve a problem of training NN for games.
 - In a single game, each consecutive pixels in a frame has high correlation from one state to another as the game environment simply does not change that rapidly in a typical Atari game.

- Such high correlation of input can cause a neural network to be highly biased and overfit with correlated experience.
- As such, we need some way to store and shuffle the input frames before feeding it forward for training of a Neural Network
- This is where Experience Replay comes in handy.
 - Experience Replay is simply a temporally Stack of replay buffer that contains of the transition information of a single action a , from one state s , to another s' . Which is denoted as follows:

$$\text{Exp Buffer} = \langle s, a, r, s' \rangle$$

- Before Starting Training Epoch
 - We store all the replay buffer to a fixed number before starting any training process.
- During Every Training Epoch
 - Random batch of experience from the replay buffer is sampled for the training process.
- After End of a Training Epoch
 - New batch of replay buffer is stored in the Experience Replay memory while the old ones are deleted

Target Network

- In the Q-learning update rule, there are two instance where a Q-function (which will be replaced by a NN) is called to calculate the target value $r + \gamma \max_a(Q(s', a'; \theta))$, and the predicted value $Q(s, a; \theta)$.
 - The problem rises when we tries to use the same Q-function for both the calculation of Q-value using the same weights from a neural network as there could be a lot of divergence between these two.
 - As small updates to Q may significantly change the policy and therefore change the data distribution, and the correlations between the action-values Q and the target values $r + \gamma \max Q(s', a')$.

An analogy will be a dog chasing a meat hang at its own tail, it will not progress forward to the desired destination but just spinning around some random spot

- To solve this problem, we make used of Target Network $Q(s', a'; \theta')$ and Predicted Network $Q(s, a; \theta)$ which is two identical neural network with different weights.

$$Loss = (r + \gamma \max_{a'} Q(s', a'; \theta') - Q(s, a; \theta))^2$$

- The target network copies the weight from the actual Q network, and is frozen from time steps to time steps.
- Freezing the target network for a while and then updating its weights with the actual Q network weights is what stabilizes the training.

A follow up analogy will be fixing a extendible pole with meat at the back of a dog and changing its length from time-to-time so that the dog will chase the meat and walk in a more stable path

Rewards Clipping

- The final tricks for training a DQN for Atari games is rewards clipping whereby the reward value is clipped to +1 to -1 across all Atari games.

Epsilon-Greedy Strategy

- Like in Q-Learning, DQN make used of Epsilon-Greedy strategy to solve the Multi-Armed Bandit problem (aka Exploration and Exploitation trade-offs) in RL. Refer to the chapter if you are unfamiliar with the concept.

The Final Algorithm

- The steps involved in a DQN are as follows:
 1. The Game Screen (state s) are first preprocod and feed to our DQN which returns the Q values of all possible actions in the state.

2. An action is selected using the epsilon-greedy policy:
 - a. With probability epsilon, we select random action a
 - b. With probability $1-\epsilon$, we select action that has maximum Q value,

$$a = \operatorname{argmax}(Q(s, a; \theta))$$
3. After selecting action a , we perform this action in state s and get the next state s' along with the reward r . Offcause s' refers to the preprocessed image of the next game screen
4. The transition is stored in our replay buffer as $\langle s, a, r, s' \rangle$
5. Next, we sample some random batches of transitions from the replay buffer and calculate the loss which is the squared difference between the target Q and predicted Q

$$Loss = (r + \gamma \max_{a'} Q(s', a'; \theta') - Q(s, a; \theta))^2$$

6. Gradient descent is then performed wrt our actual network parameters θ to minimize the loss
7. After every k steps, we copy the actual network weights θ to our target network weights, θ'
8. We repeat these steps for M number of episodes.

Extra Tips and Tricks

Double DQN

Prioritized Experience Replay

Dueling Network Architecture